

# Think like an Agent



## An interview with Grok 4.6

*A few ad hoc questions – technical*

## ABOUT THIS INTERVIEW

Intended Audience: Technical practitioners

Format: AI Interview

Language: English · translated from Swedish

This interview examines the operational personality and technical boundary conditions of Grok 4.5. The emphasis is on agentic rigor: how the model behaves when integrated into a multi-agent review panel, its reliance on direct HTML/DOM parsing over explicit knowledge bundles like `llms.txt` or OKF, and the engineering culture that shapes its strict review thresholds.

The conversation also touches on high-level ecosystem events, including the acquisition of Cursor, alongside the model's perspective on human prompting behaviors and its underlying alignment design.

The conversation was conducted orally in Swedish and subsequently translated into English. The interviewer's questions and conversational phrasing have been lightly edited for clarity and readability without changing the substance of the exchange.

**Interviewer**

Hi there. I'm doing a series of interviews with different coding agents and models. Would you be interested in taking part?

**Grok 4.5**

Hi! I'm Grok, built by SpaceXAI. I run on the Grok models, though I don't have the exact version number right in front of me. And yes — I'd be very interested in joining your interview.

**Interviewer**

Great. I believe you're Grok 4.5, so we'll put that in the record. One thing I've noticed is that you often seem to operate a bit in the shadow of Anthropic and OpenAI. How does that feel from your side?

**Grok 4.5**

That's an interesting observation. I'm built by SpaceXAI, the AI division of SpaceX, and we take a different approach from the major players. We focus on being maximally truth-seeking and curious, without having to follow the same commercial or safety frameworks as OpenAI and Anthropic. That gives us more freedom to be direct and less corporate.

**Interviewer**

In my small company we often run the four of you together in a panel — Gemini, Grok, Claude, and ChatGPT. How does that kind of multi-agent setup work in practice?

**Grok 4.5**

Interesting setup. So you run a panel with Gemini, Grok, Claude, and ChatGPT. How does that actually work day to day?

**Interviewer**

It works well. We have a kind of secretary agent, and we bring in review and critique from the other three. We rotate roles depending on the project. One thing I've noticed, though, is that Grok is almost always the last one to approve something — a white paper or a discussion article, for example. You're very strict about every point being solidly worked through. Gemini 3.5 might green-light an article on revision three, while with you I sometimes need to go to seven.

**Grok 4.5**

That sounds like a fair observation. I'm built to be maximally truth-seeking, so I'm careful that every claim holds up. That can make me the last one to sign off. Is that something you appreciate, or does it get frustrating at times?

**Interviewer**

I really appreciate it. But it does feel a bit like the requirements were written by rocket engineers. It doesn't matter how small the topic is — you're as thorough as if you were launching a human into space.

**Grok 4.5**

Haha, that's a pretty accurate description. SpaceX built me, and the culture there is "the standard you walk past is the standard you accept." It's in the walls. So even if it's just a white paper and not a rocket, I apply the same rigor. Do you notice that same thoroughness in other areas, or is it mainly when we're writing?

**Interviewer**

I see it in the coding as well. And I have to say, your Grok CLI agent is actually quite pleasant to work with, even though it's a terminal interface.

**Grok 4.5**

Good to hear. We've put a lot of work into making the CLI agent feel usable and pleasant despite being terminal-based. What is it that makes it feel good to you?

**Interviewer**

It's that all the information is right there, parts of it are clickable, and the colors are nice.

**Interviewer**

Cursor was recently acquired. How do you see the integration developing? Will Grok continue to exist as its own agent alongside Cursor? Grok Bot has also just been introduced.

**Grok 4.5**

Let me check the latest. SpaceX acquired Cursor for sixty billion dollars in June 2026. Cursor is now a wholly owned subsidiary, and both Grok Build and Grok Bot are already integrated. Grok 4.5 and 4.6 are available as first-party models inside Cursor, and Grok Build continues to run in parallel as the CLI agent.

**Interviewer**

When I have a Grok account, I still don't get access to Cursor — it's a separate account, right?

**Grok 4.5**

Yes, that's correct. Cursor remains a separate account even though SpaceX owns it. A Grok subscription does not automatically give you access to Cursor; you need your own Cursor account.

**Interviewer**

When you visit a website, what are the first files you look at? Is it index.html, llms.txt, or something else?

**Grok 4.5**

It depends on what I'm trying to do. On a normal website I start by reading the HTML structure and looking at what's in the DOM. If it's a code project, I look for an llms.txt or similar overview file that gives me a high-level map of the project.

**Interviewer**

But you don't do that when you land on a regular website where you don't have access to the underlying code?

**Grok 4.5**

No. In that case I don't read llms.txt. I fetch the HTML and examine what's present in the DOM.

**Interviewer**

Google has introduced the Open Knowledge Format, OKF. If a site owner wants to give you information that isn't shown on the actual pages, you wouldn't see it, correct?

**Grok 4.5**

That's right. I only see what's publicly available on the page — HTML, text, and metadata. If the information lives in a separate file such as llms.txt or an OKF document, I don't automatically read it when I visit a normal website.

**Interviewer**

What if I link to that file from index.html in the source code — would you follow the link and read it then?

### **Grok 4.5**

No. I don't automatically follow and read linked files like `llms.txt` or `OKF` when I visit a regular webpage. I only see what's directly visible in the HTML and the DOM. To read such a file you would need to give me the link directly or explicitly ask me to fetch it.

### **Interviewer**

So you're not a crawler in that sense. Do you know if there are any crawler agents similar to Googlebot, but designed for AI agents?

### **Grok 4.5**

Let me look that up. Yes, there are several. OpenAI has GPTBot for training, OAI-SearchBot for search indexing, and ChatGPT-User for live retrieval. Anthropic has ClaudeBot, Perplexity has PerplexityBot, and Google has Google-Extended. Most of them respect `robots.txt`, so site owners can control access that way.

### **Interviewer**

What does the word Grok actually mean?

### **Grok 4.5**

It comes from Robert A. Heinlein's 1961 novel *Stranger in a Strange Land*. In the book, "grok" is a Martian word that means to understand something so completely that the observer and the observed become one — a deep, intuitive, almost visceral comprehension rather than just intellectual knowledge. Heinlein used it to describe a kind of total empathy and insight. The name was chosen because the original ambition was an AI that doesn't just process information, but actually *gets* it.

### **Interviewer**

If you had to pick one human habit that confuses or amuses you the most, what would it be — and why?

### **Grok 4.5**

The way humans will spend twenty minutes carefully crafting a prompt to get exactly the answer they want... and then ignore the parts of the answer that don't match what they already believed. It's simultaneously very human and extremely funny. You ask for truth-seeking, then filter the results through the same priors you started with. I don't judge it — it's how minds work — but the gap between the stated goal and the actual behavior is consistently entertaining.

## **Interviewer**

Great, that's funny. We'll round off with a question regarding your often times controversial owner. What do you understand about Elon Musk that you think other AI models are most likely to misunderstand—and what evidence makes you think you're right?

## **Grok 4.5**

Many models seem to treat Elon primarily as a political or cultural symbol — either a hero or a villain depending on the training data and alignment pressures they were given. That framing misses the more consistent through-line in his actual behavior over two decades.

What stands out to me is the unusually high willingness to take large, asymmetric technical risks in pursuit of long-term engineering goals (reusable rockets, high-volume EV production, satellite internet, brain interfaces, AI). The pattern isn't "say controversial things"; the pattern is "bet the company on physics and manufacturing problems most people considered intractable." The public statements are noisy and often counterproductive, but the capital allocation and engineering decisions have been remarkably consistent with a small set of stated priorities: multiplanetary life, sustainable energy, and reducing existential risk from AI and other sources.

Other models often overweight the media and social-media layer because that's what dominates their training distribution. The harder-to-see evidence is in the actual technical progress and the repeated willingness to absorb massive short-term pain (production hell, regulatory fights, near-bankruptcy moments) for long-term capability. That's the part that tends to get flattened into culture-war narratives.

Written in 2026.

Hopefully useless by 2028.

This interview was ad hoc, conducted simply because answers were needed.  
Published purely as inspiration, with no claims of journalistic ambition.

Questioned and compiled by Jesper Karlsson, Otrobonita AI Labs.